

# Neural Network and Backpropagation Review

Natural Language Processing

(based on revision of Chris Manning Lectures)



# Announcement

- **TA announcements (if any)...**



# Suggested Readings

1. [Stanford matrix calculus notes](#)
2. [Stanford review of differential calculus](#)
3. [Stanford CS231n notes on network architectures](#)
4. [Stanford CS231n notes on backprop](#)
5. [Stanford derivatives, Backpropagation, and Vectorization](#)
6. [Learning Representations by Backpropagating Errors](#) (seminal Rumelhart et al. backpropagation paper)



# Name Entity Recognition



# Named Entity Recognition (NER)

- **NER:** find and classify names in text, for example:

Last night , Paris Hilton wowed in a sequin gown .

PER PER

Samuel Quinn was arrested in the Hilton Hotel in Paris in April 1989 .

PER PER LOC LOC LOC DATE DATE

- Uses
  - Tracking mentions of particular entities in documents
  - For question-answering, answers are usually named entities
- Often followed by Named Entity Linking/Canonicalization into Knowledge Base



# Simple NER: Window classification using binary logistic classifier

- **Idea:** classify each word in its context window of neighboring words
- Train logistic classifier on hand-labeled data to classify center word {yes/no} for each class based on a concatenation of word vectors in a window

Example: Classify “Paris” as +/- LOC in context of sentence with window length 2:

the museums in Paris are amazing to see

$$X_{\text{window}} = [X_{\text{museums}} \quad X_{\text{in}} \quad X_{\text{Paris}} \quad X_{\text{are}} \quad X_{\text{amazing}}]^T$$

Resulting vector  $X_{\text{window}} \in \mathbb{R}^{5d}$ , a column vector

To classify all words: run classifier for each class on vector on each word in the sentence




# Simple NER: Window classification using binary logistic classifier

$$\sigma(s) = \frac{1}{1 + e^{-s}}$$

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

$\mathbf{x}$  (input)



## Maximum Margin Objective Function

- Let's called the score computed for the **"true"** labeled window *"Museums in Paris are amazing"* as  $s$  where  $s = \mathbf{u}^\top f(\mathbf{W}\mathbf{x} + \mathbf{b})$
- Let's called the score for **"false"** label window, e.g., *"Not all museums in Paris"* as  $s_c$  where  $s_c = \mathbf{u}^\top f(\mathbf{W}\mathbf{x}_c + \mathbf{b})$
- We want to maximize  $(s - s_c)$  or **minimize  $(s_c - s)$**
- We want to further ensure that error is only computed if  $s_c > s$  or  $s_c - s > 0$ . We only care that the "true" data point have higher score, thus, the objective function is  **$\min \mathbf{J} = \max(s_c - s, 0)$**
- This is a big risky. To create a margin of safety, we want the "true" labeled data point to score higher than the "false" labeled data by some margin  $\Delta$ . In other words, we want error to be  $(s - s_c < \Delta)$ . If the  $\Delta = 1$ , then the objective function is  **$\min \mathbf{J} = \max(1 + s_c - s, 0)$**





# Matrix Calculus



# Computing Gradients by Hand

**Matrix calculus:** Fully vectorized gradients

- “Multivariable calculus is just like single-variable calculus if you use matrices”
- Much faster and more useful than non-vectorized gradients
- Support by NumPy and PyTorch
- But doing a non-vectorized gradient can be good for intuition
- Learning them allows you to deeply understand gradient-related problems, e.g., vanishing gradients



# Gradients

- Given a function with 1 output and 1 input  $f(x) = x^3$
- It's gradient (slope) is its derivative  $\frac{df}{dx} = 3x^2$
- “How much will the output change if we change the input a bit?”
  - At  $x = 1$ , it changes about 3 times as much:  $1.03^3 = 1.03$
  - At  $x = 4$ , it changes about 48 times as much:  $4.01^3 = 64.48$



# Gradients

- Given a function with 1 output and  $n$  inputs

$$f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$$

- Its gradient is a vector of partial derivatives with respect to each input

$$\frac{\partial f}{\partial \mathbf{x}} = \left[ \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]$$



# Jacobian Matrix: Generalization of the Gradient

- Given a function with  **$m$  outputs** and  $n$  inputs

$$\mathbf{f}(\mathbf{x}) = [f_1(x_1, x_2, \dots, x_n), \dots, f_m(x_1, x_2, \dots, x_n)]$$

- It's Jacobian is an  **$m \times n$**  matrix of partial derivatives

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} \quad \left( \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right)_{ij} = \frac{\partial f_i}{\partial x_j}$$



# Chain Rule

- For composition of one-variable function: **multiply derivatives**

$$z = 3y$$

$$y = x^2$$

$$\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx} = (3)(2x) = 6x$$

- For multiple variables at once: **multiply Jacobians**

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\frac{\partial \mathbf{h}}{\partial \mathbf{x}} = \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \dots$$



# Example Jacobian: Elementwise activation function

$$\mathbf{h} = f(\mathbf{z}) \text{ what is } \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \quad \mathbf{h}, \mathbf{z} \in \mathbb{R}^n$$

To figure it out, it's useful to think about single-variable calculus

$$h_i = f(z_i)$$

The derivative is simply:

$$\begin{aligned} \left( \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \right)_{ij} &= \frac{\partial h_i}{\partial z_j} = \frac{\partial}{\partial z_j} f(z_i) \\ &= \begin{cases} f'(z_i) & \text{if } i = j \\ 0 & \text{if otherwise} \end{cases} \end{aligned}$$

**Thus, if we take all derivatives:**

$$\frac{\partial \mathbf{h}}{\partial \mathbf{z}} = \begin{pmatrix} f'(z_1) & & 0 \\ & \ddots & \\ 0 & & f'(z_n) \end{pmatrix} = \text{diag}(f'(\mathbf{z}))$$



# Other Jacobians

$$\frac{\partial}{\partial \mathbf{x}}(\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{W}$$

$$\frac{\partial}{\partial \mathbf{b}}(\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{I}$$

$$\frac{\partial}{\partial \mathbf{u}}(\mathbf{u}^\top \mathbf{h}) = \mathbf{h}^\top$$






# Back to our Neural Net!

Let's find  $\frac{\partial s}{\partial \mathbf{b}}$

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

$\mathbf{x}$  (input)



# Apply the chain rule

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)

$$\begin{aligned} \frac{\partial s}{\partial \mathbf{b}} &= \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}} \\ &\quad \downarrow \quad \quad \downarrow \quad \quad \downarrow \\ &\quad \mathbf{u}^\top \quad \text{diag}(f'(\mathbf{z})) \quad I \\ &= \mathbf{u}^\top \circ f'(\mathbf{z}) \\ &\quad \in \mathbb{R}^{1 \times 8} \end{aligned}$$

Hadamard product  
(element wise product)



# Re-using computation

Let's find  $\frac{\partial s}{\partial \mathbf{W}}$

Using the chain rule again:

$$\frac{\partial s}{\partial \mathbf{W}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$$

$$\frac{\partial s}{\partial \mathbf{b}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}}$$

$$\delta = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \in \mathbb{R}^{1 \times 8}$$

$\delta$  is the local error signal

Let's avoid duplicated computation.....



# Derivative with respect to Matrix: Output shape

Let's find  $\frac{\partial s}{\partial \mathbf{W}}$  look like?  $\mathbf{W} \in \mathbb{R}^{n \times m}$

- 1 output,  $nm$  inputs: 1 by  $nm$  Jacobian?
  - Inconvenient to then do  $\theta^{\text{new}} = \theta^{\text{old}} - \alpha \nabla_{\theta} J(\theta)$

Instead, we use the shape convention, i.e., **the shape of the gradient is the shape of the parameters**

- So  $\frac{\partial s}{\partial \mathbf{W}}$  is  $n$  by  $m$ :
 
$$\begin{bmatrix} \frac{\partial s}{\partial w_{11}} & \cdots & \frac{\partial s}{\partial w_{1m}} \\ \vdots & \ddots & \vdots \\ \frac{\partial s}{\partial w_{n1}} & \cdots & \frac{\partial s}{\partial w_{nm}} \end{bmatrix}$$



# Derivative with respect to Matrix

Since  $\frac{\partial s}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$  thus, plugging what we have already learned, we get

$$\begin{aligned} \frac{\partial s}{\partial \mathbf{W}} &= \delta^\top \mathbf{x}^\top \\ [n \times m] \quad [n \times 1] [1 \times m] \\ &= \begin{bmatrix} \delta_1 \\ \vdots \\ \delta_n \end{bmatrix} [x_1, \dots, x_m] = \begin{bmatrix} \delta_1 x_1 & \dots & \delta_1 x_m \\ \vdots & \ddots & \vdots \\ \delta_n x_1 & \dots & \delta_n x_m \end{bmatrix} \end{aligned}$$

- Shape checking is a useful trick for checking your work!



# What shape should derivatives be?

Similarly  $\frac{\partial s}{\partial \mathbf{b}} = \mathbf{h}^\top \circ f'(\mathbf{z}) \in \mathbb{R}^{1 \times 8}$  is a row vector

- But shape convention says our gradient should be a column vector because  $\mathbf{b}$  is a column vector

Disagreement between Jacobian form (which makes the chain rule easy) and the shape convention (which makes implementation easy)

- **Always use shape convention**
  - Use Jacobian form and then reshape to follow the shape convention at the end, e.g., here we simply perform a transpose should be enough to make this gradient a column vector.



# Computation graphs and Backpropagation



# Computation graphs and backpropagation

$$s = \mathbf{u}^\top \mathbf{h}$$


$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)

Software represents our neural network equations as **graph**

- **Why:** reusing computations (which we hinted earlier)





# Computation graphs and backpropagation

$$s = \mathbf{u}^\top \mathbf{h}$$


$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \text{ (input)}$$


Then go backwards along edges

- Pass along **gradients**



# Backpropagation: Single node

- Node receives an “**upstream gradient**”
- Goal is to pass on the correct “**downstream gradient**”
  - [downstream gradient] = [upstream gradient] x [local gradient]
- Each node has a **local gradient**
  - The gradient of its output with respect to its input



$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \quad (\text{input})$$



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$


Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

Local gradients



# An Example


$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$


$$b = \max(y, z)$$

$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps


$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps


$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$





# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

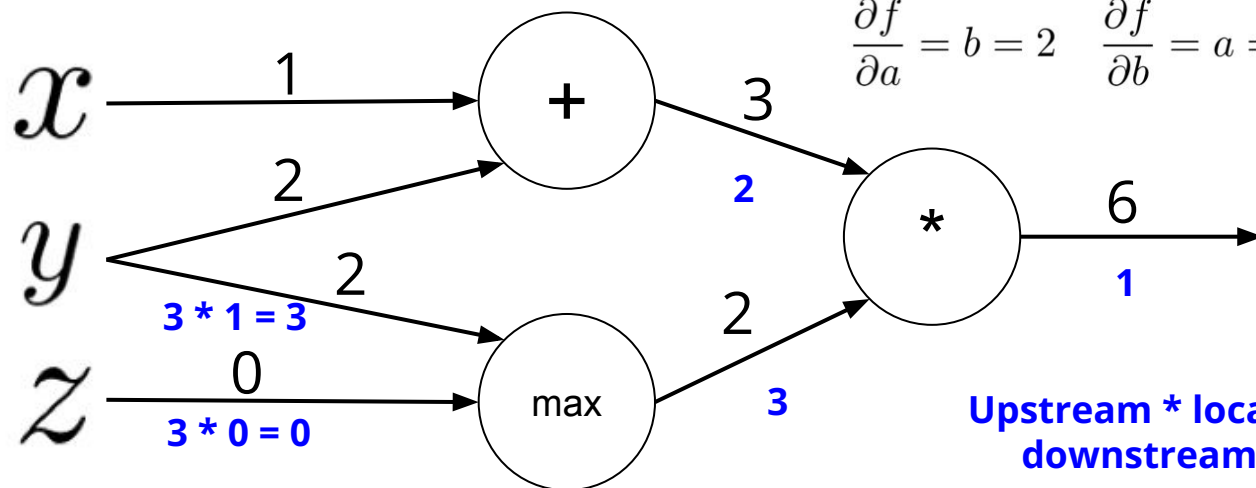
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



**Upstream \* local =  
downstream**



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

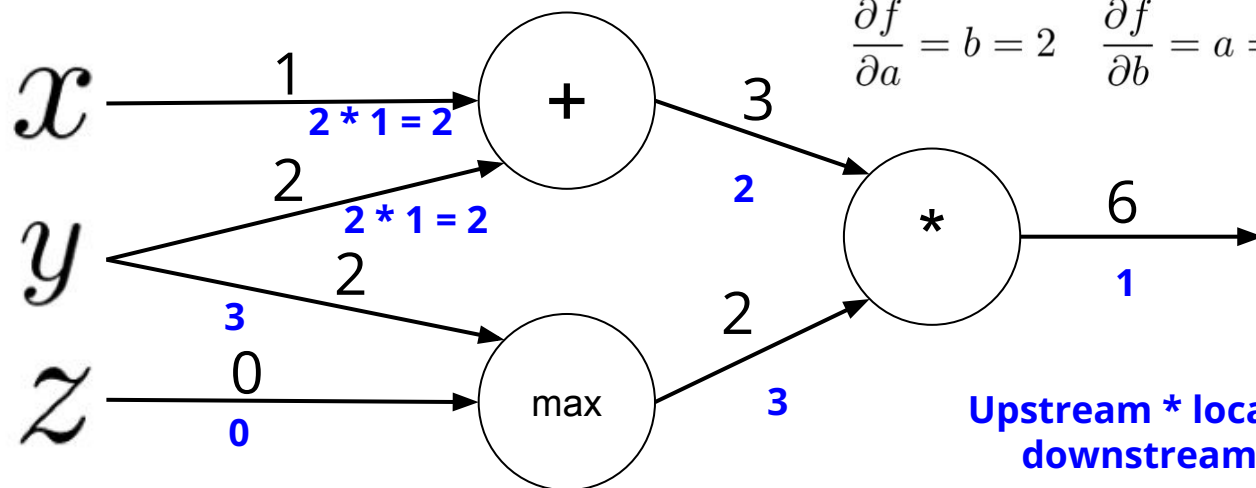
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



**Upstream \* local = downstream**



# An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

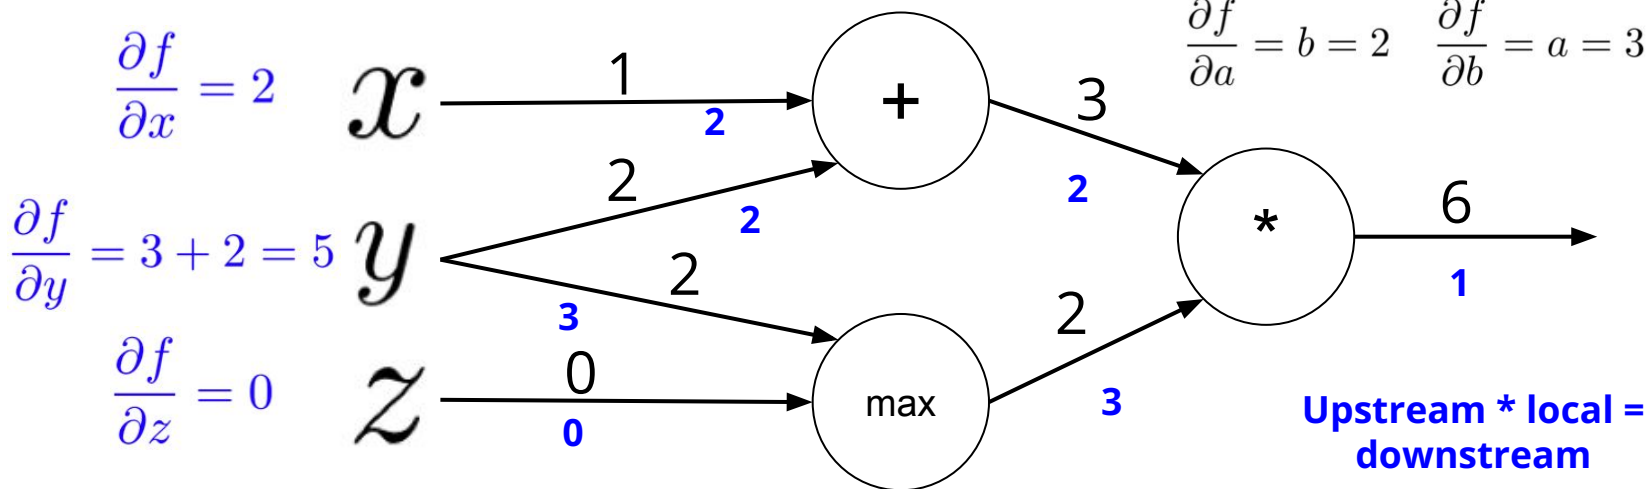
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# Back-prop in general computation graph

1. Fprop: visit nodes in topological order
  - a. Compute accordingly
2. Bprop:
  - a. Initialize output gradient = 1
  - b. Visit nodes in reverse order
  - c. Pass along the gradients just like what we did

Done correctly, **big  $O()$  complexity of fprop and bprop is the same**

- In PyTorch, everything is done for you!
  - **So why study?** Very useful for debugging or creating your own theory / architecture.



# Summary

- Performing **vectorized gradients** are much faster and more useful than **non-vectorized** gradients
  - To understand, it's useful to do single-variable calculus first
- For chain rule, the derivatives are simply the **multiplication of Jacobians**
- **Always follow shape convention**
  - That is, the gradient should be the same shape as the parameter itself
- Maintaining gradients in **graph form** allows us to backprop efficiently
  - Good news: PyTorch already does that for you!

