

Neural Network and Backpropagation Review

Natural Language Processing

(based on revision of Chris Manning Lectures)



Announcement

- TA announcements (if any)...



Suggested Readings

1. [Stanford matrix calculus notes](#)
2. [Stanford review of differential calculus](#)
3. [Stanford CS231n notes on network architectures](#)
4. [Stanford CS231n notes on backprop](#)
5. [Stanford derivatives, Backpropagation, and Vectorization](#)
6. [Learning Representations by Backpropagating Errors](#) (seminal Rumelhart et al. backpropagation paper)



Name Entity Recognition



Named Entity Recognition (NER)

- **NER:** find and classify names in text, for example:

Last night , Paris Hilton wowed in a sequin gown .

PER PER

Samuel Quinn was arrested in the Hilton Hotel in Paris in April 1989 .

PER PER LOC LOC LOC DATE DATE

- Uses
 - Tracking mentions of particular entities in documents
 - For question-answering, answers are usually named entities
- Often followed by Named Entity Linking/Canonicalization into Knowledge Base



Simple NER: Window classification using binary logistic classifier

- **Idea:** classify each word in its context window of neighboring words
- Train logistic classifier on hand-labeled data to classify center word {yes/no} for each class based on a concatenation of word vectors in a window

Example: Classify "Paris" as +/- LOC in context of sentence with window length 2:

$$\begin{array}{ccccccccc} \text{the} & & \text{museums} & \text{in} & \text{Paris} & \text{are} & \text{amazing} & \text{to} & \text{see} \\ X_{\text{window}} = [X_{\text{museums}} & X_{\text{in}} & X_{\text{Paris}} & X_{\text{are}} & X_{\text{amazing}}]^T \end{array}$$

Resulting vector $X_{\text{window}} \in \mathbb{R}^{5d}$, a column vector

To classify all words: run classifier for each class on vector on each word in the sentence




Simple NER: Window classification using binary logistic classifier

$$\sigma(s) = \frac{1}{1 + e^{-s}}$$

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

\mathbf{x} (input)



Maximum Margin Objective Function

- Let's called the score computed for the **"true"** labeled window *"Museums in Paris are amazing"* as s where $s = \mathbf{u}^\top f(\mathbf{W}\mathbf{x} + \mathbf{b})$
- Let's called the score for **"false"** label window, e.g., *"Not all museums in Paris"* as s_c where $s_c = \mathbf{u}^\top f(\mathbf{W}\mathbf{x}_c + \mathbf{b})$
- We want to maximize $(s - s_c)$ or **minimize $(s_c - s)$**
- We want to further ensure that error is only computed if $s_c > s$ or $s_c - s > 0$. We only care that the "true" data point have higher score, thus, the objective function is **$\min \mathbf{J} = \max(s_c - s, 0)$**
- This is a big risky. To create a margin of safety, we want the "true" labeled data point to score higher than the "false" labeled data by some margin Δ . In other words, we want error to be $(s - s_c < \Delta)$. If the $\Delta = 1$, then the objective function is **$\min \mathbf{J} = \max(1 + s_c - s, 0)$**



Matrix Calculus



Computing Gradients by Hand

Matrix calculus: Fully vectorized gradients

- “Multivariable calculus is just like single-variable calculus if you use matrices”
- Much faster and more useful than non-vectorized gradients
- Support by NumPy and PyTorch
- But doing a non-vectorized gradient can be good for intuition
- Learning them allows you to deeply understand gradient-related problems, e.g., vanishing gradients



Gradients

- Given a function with 1 output and 1 input $f(x) = x^3$
- It's gradient (slope) is its derivative $\frac{df}{dx} = 3x^2$
- *"How much will the output change if we change the input a bit?"*
 - At $x = 1$, it changes about 3 times as much: $1.03^3 = 1.03$
 - At $x = 4$, it changes about 48 times as much: $4.01^3 = 64.48$



Gradients

- Given a function with 1 output and n inputs

$$f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$$

- Its gradient is a vector of partial derivatives with respect to each input

$$\frac{\partial f}{\partial \mathbf{x}} = \left[\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]$$



Jacobian Matrix: Generalization of the Gradient

- Given a function with **m outputs** and n inputs

$$\mathbf{f}(\mathbf{x}) = [f_1(x_1, x_2, \dots, x_n), \dots, f_m(x_1, x_2, \dots, x_n)]$$

- It's Jacobian is an **$m \times n$** matrix of partial derivatives

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \dots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} \quad \left(\frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right)_{ij} = \frac{\partial f_i}{\partial x_j}$$



Chain Rule

- For composition of one-variable function: **multiply derivatives**

$$z = 3y \quad y = x^2 \quad \frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx} = (3)(2x) = 6x$$

- For multiple variables at once: **multiply Jacobians**

$$\mathbf{h} = f(\mathbf{z}) \quad \mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad \frac{\partial \mathbf{h}}{\partial \mathbf{x}} = \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \dots$$



Example Jacobian: Elementwise activation function

$$\mathbf{h} = f(\mathbf{z}) \text{ what is } \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \quad \mathbf{h}, \mathbf{z} \in \mathbb{R}^n$$

To figure it out, it's useful to think about single-variable calculus

$$h_i = f(z_i)$$

The derivative is simply:

$$\begin{aligned} \left(\frac{\partial \mathbf{h}}{\partial \mathbf{z}} \right)_{ij} &= \frac{\partial h_i}{\partial z_j} = \frac{\partial}{\partial z_j} f(z_i) \\ &= \begin{cases} f'(z_i) & \text{if } i = j \\ 0 & \text{if otherwise} \end{cases} \end{aligned}$$

Thus, if we take all derivatives:

$$\frac{\partial \mathbf{h}}{\partial \mathbf{z}} = \begin{pmatrix} f'(z_1) & & 0 \\ & \ddots & \\ 0 & & f'(z_n) \end{pmatrix} = \text{diag}(f'(\mathbf{z}))$$



Other Jacobians

$$\frac{\partial}{\partial \mathbf{x}}(\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{W}$$

$$\frac{\partial}{\partial \mathbf{b}}(\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{I}$$

$$\frac{\partial}{\partial \mathbf{u}}(\mathbf{u}^\top \mathbf{h}) = \mathbf{h}^\top$$




Back to our Neural Net!

Let's find $\frac{\partial s}{\partial \mathbf{b}}$

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

\mathbf{x} (input)



Apply the chain rule

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

\mathbf{x} (input)

$$\begin{aligned} \frac{\partial s}{\partial \mathbf{b}} &= \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}} \\ &\quad \downarrow \quad \quad \downarrow \quad \quad \downarrow \\ &\quad \mathbf{u}^\top \quad \text{diag}(f'(\mathbf{z})) \quad I \\ &= \mathbf{u}^\top \circ f'(\mathbf{z}) \\ &\quad \in \mathbb{R}^{1 \times 8} \end{aligned}$$

Hadamard product
(element wise product)



Re-using computation

Let's find $\frac{\partial s}{\partial \mathbf{W}}$

Using the chain rule again:

$$\frac{\partial s}{\partial \mathbf{W}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$$

$$\frac{\partial s}{\partial \mathbf{b}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}}$$

$$\delta = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \in \mathbb{R}^{1 \times 8}$$

δ is the local error signal

Let's avoid duplicated computation.....



Derivative with respect to Matrix: Output shape

Let's find $\frac{\partial s}{\partial \mathbf{W}}$ look like? $\mathbf{W} \in \mathbb{R}^{n \times m}$

- 1 output, nm inputs: 1 by nm Jacobian?
 - Inconvenient to perform gradient update

$$\theta^{\text{new}} = \theta^{\text{old}} - \alpha \nabla_{\theta} J(\theta)$$

Instead, we use the shape convention, i.e., **the shape of the gradient is the shape of the parameters**

- So $\frac{\partial s}{\partial \mathbf{W}}$ is n by m :

$$\begin{bmatrix} \frac{\partial s}{\partial w_{11}} & \cdots & \frac{\partial s}{\partial w_{1m}} \\ \vdots & \ddots & \vdots \\ \frac{\partial s}{\partial w_{n1}} & \cdots & \frac{\partial s}{\partial w_{nm}} \end{bmatrix}$$



Derivative with respect to Matrix

Since $\frac{\partial s}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$ thus, plugging what we have already learned, we get

$$\begin{aligned} \frac{\partial s}{\partial \mathbf{W}} &= \delta^\top \mathbf{x}^\top \\ [n \times m] \quad [n \times 1][1 \times m] \\ &= \begin{bmatrix} \delta_1 \\ \vdots \\ \delta_n \end{bmatrix} [x_1, \dots, x_m] = \begin{bmatrix} \delta_1 x_1 & \dots & \delta_1 x_m \\ \vdots & \ddots & \vdots \\ \delta_n x_1 & \dots & \delta_n x_m \end{bmatrix} \end{aligned}$$

Shape checking is a useful trick for checking your work!



What shape should derivatives be?

Similarly $\frac{\partial s}{\partial \mathbf{b}} = \mathbf{h}^\top \circ f'(\mathbf{z}) \in \mathbb{R}^{1 \times 8}$ is a row vector

- But shape convention says our gradient should be a column vector because b is a column vector

Disagreement between Jacobian form (which makes the chain rule easy) and the shape convention (which makes implementation easy)

- **Always use shape convention**
 - Use Jacobian form and then reshape to follow the shape convention at the end, e.g., here we simply perform a transpose should be enough to make this gradient a column vector.



Computation graphs and Backpropagation



Computation graphs and backpropagation

$$s = \mathbf{u}^\top \mathbf{h}$$


$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \text{ (input)}$$

Software represents our neural network equations as **graph**

- **Why:** reusing computations (which we hinted earlier)



Computation graphs and backpropagation

Then go backwards along edges


- Pass along **gradients**

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \text{ (input)}$$



Backpropagation: Single node


- Node receives an **“upstream gradient”**
- Goal is to pass on the correct **“downstream gradient”**
 - [downstream gradient] = [upstream gradient] x [local gradient]
- Each node has a **local gradient**
 - The gradient of its output with respect to its input

$$s = \mathbf{u}^\top \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

\mathbf{x} (input)




An Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients



An Example


$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$


Forward prop steps

$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps


$$a = x + y$$
$$b = \max(y, z)$$
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

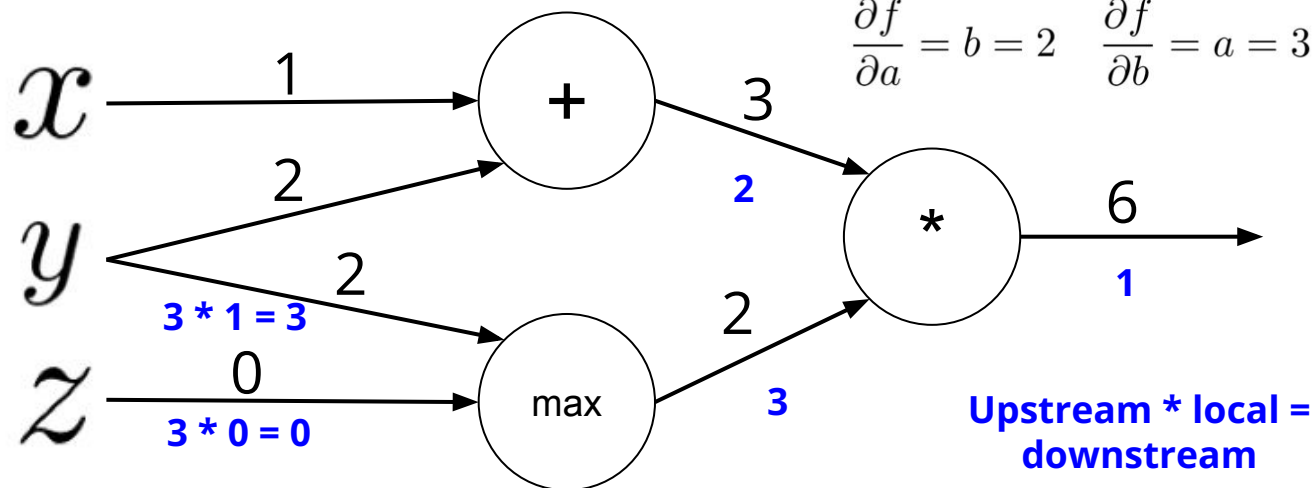
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

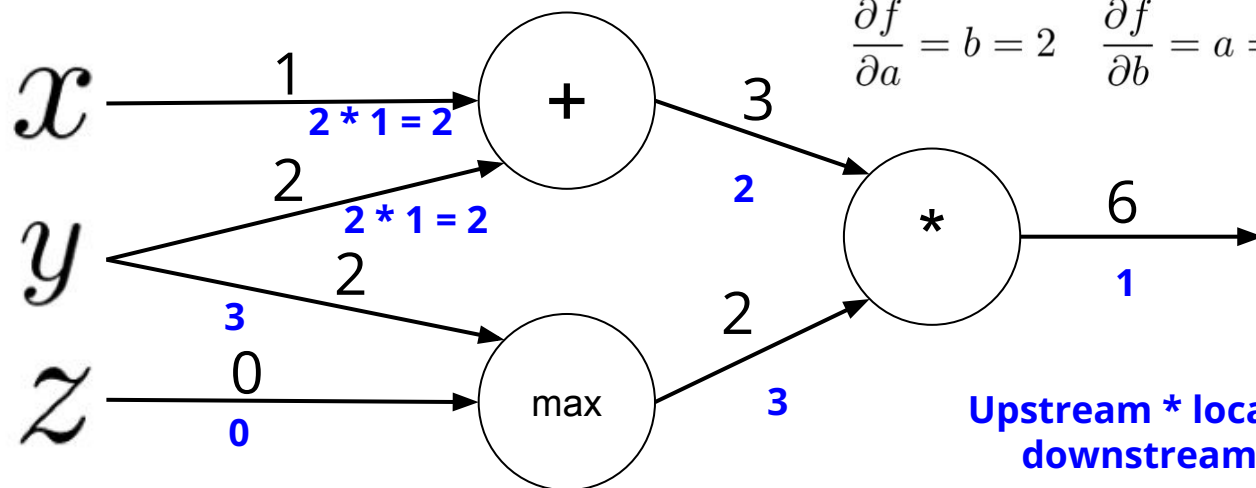
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



**Upstream * local =
downstream**



An Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

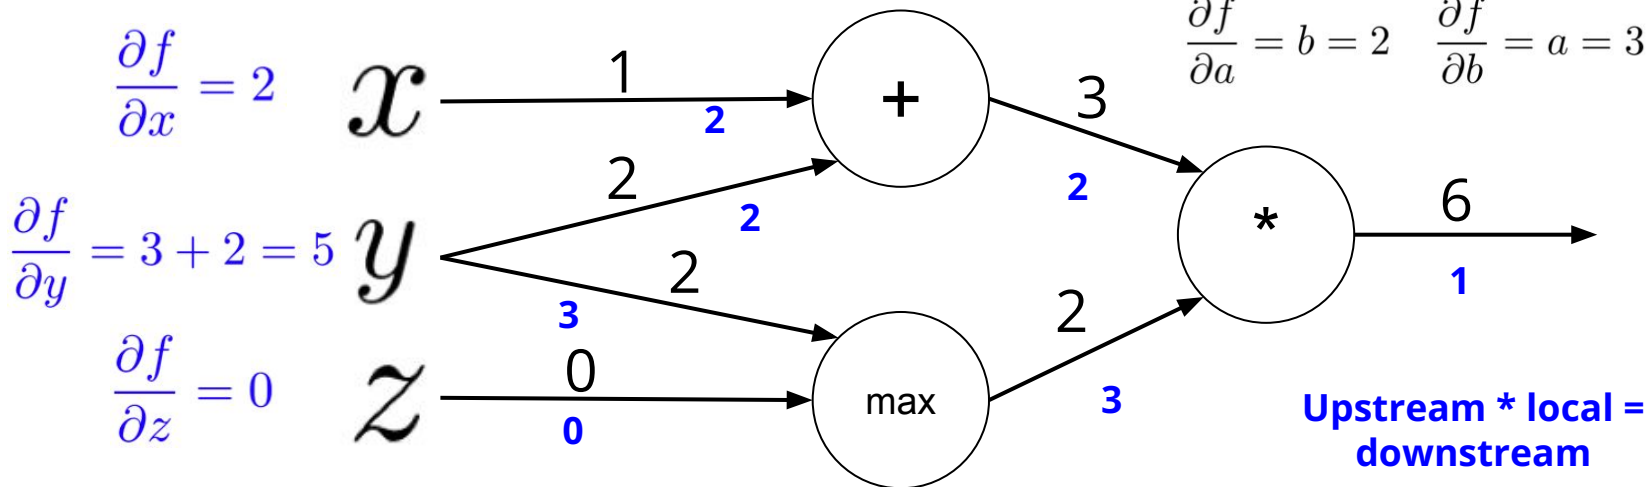
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



Back-prop in general computation graph

1. Fprop: visit nodes in topological order
 - a. Compute accordingly
2. Bprop:
 - a. Initialize output gradient = 1
 - b. Visit nodes in reverse order
 - c. Pass along the gradients just like what we did

Done correctly, **big O() complexity of fprop and bprop is the same**

- In PyTorch, everything is done for you!
 - **So why study?** Very useful for debugging or model development



Summary

- Performing **vectorized gradients** are much faster and more useful than **non-vectorized** gradients
 - To understand, it's useful to do single-variable calculus first
- For chain rule, the derivatives are simply the **multiplication of Jacobians**
- **Always follow shape convention**
 - That is, the gradient should be the same shape as the parameter itself
- Maintaining gradients in **graph form** allows us to backprop efficiently
 - Good news: PyTorch already does that for you!

